

Efficient Numerical Code in Haskell

Jyotirmoy Bhattacharya

Ambedkar University Delhi

Delhi NCR Haskell Meetup, June 2016

About

- Email: jyotirmoy@jyotirmoy.net
- Website: <https://www.jyotirmoy.net>
- GitHub: <https://github.com/jmoy>

Lists: The Good

The List datatype

```
data List a = Empty  
            | Node a (List a)
```

- Lazy links can represent infinite and self-referential structures:
`fib = 1:1:zipWith (+) fib (tail fib)`
- Lazy data can support modularity:
`filter (even . fst) $ map (\i->(i,i^2)) [1..]`
- Updates can share representation with old values:

```
x = "cake"  
y = 'b':tail x
```

Lists: The Bad

- Access to the n -th elements in $O(n)$.
- High memory overhead and poor data locality.
- Pointer chasing for traversal.
- Data and links must be forced before use.

Sieve of Eratosthenes (Python)

```
import numpy as np
def count_primes(n):
    "Count no. of primes smaller than n"
    v = np.full(n, True, np.bool)
    # We want to set v[i] = False for non-primes
    v[0] = v[1] = False

    for p in range(n):
        if v[p]:
            # Found a prime
            # Mark all its multiples as non-prime
            for k in range(p*p, n, p):
                v[k] = False
    return np.count_nonzero(v)
```

The vector package

- Module `Data.Vector`. Array of heap-allocated values.
- Module `Data.Vector.Unboxed`. Array of unboxed values.
- Module `Data.Vector.Storable`. Array allocated in memory that can be passed to foreign functions.

Sieve of Eratosthenes (Haskell)

```
import qualified Data.Vector.Unboxed as V
import Data.Vector.Unboxed ((//),(!))
countPrimes n = V.length $ V.filter id $ loop 2 mask0
where
    mask0 = (V.replicate n True) // [(0,False),(1,False)]
    loop p mask
        | p == n = mask
        | not (mask ! p) = loop (p+1) mask
        | otherwise = loop (p+1) mask'
            where
                mask' = mask // [(k,False)|k<-[p*p,p*p+p..(n-1)]]
```

The ST Monad

- A value of type $(\text{ST } s \ a)$ is a computation which transforms a state of type `s` and additionally produces a value of type `a`.
- The state is not accessible to library users. So they cannot duplicate it.
- Action `newSTRef` creates a new *mutable variable* within a state to which values can be written through `writeSTRef` and `modifySTRef` and read through `readSTRef`.
- The `runST` function takes a computation in `ST` and gives us the value produced. This is what allows us to embed `ST` computation in pure programs.

Stupid ST example

```
import Control.Monad.ST
import Data.STRef
import Control.Monad (forM_)

lengthJB :: [a] -> Int
lengthJB xs = runST $ do
    ctr <- newSTRef 0
    forM_ xs $ \_ -> modifySTRef ctr (+ 1)
    readSTRef ctr
```

Simulating a bank

The External Interface

```
type Name = String
type Ledger = [(Name,Double)]
data Transaction = Transaction
    Name      -- From
    Name      -- To
    Double   -- Amount
```

-- How do we implement this efficiently?

```
simulateBank :: Ledger -> [Transaction] -> Ledger
```

Simulating a bank (contd.)

```
-- The state of a bank
type Bank s = HashMap Name (STRef s Double)
-- Create a bank reflecting a given ledger
mkBank :: [(Name,Double)] -> ST s (Bank s)
-- Modify the state of a bank to reflect a transaction
transact :: Bank s -> Transaction -> ST s ()
-- Turn the state of the bank into a ledger
readLedger :: Bank s -> ST s Ledger

simulateBank :: Ledger -> [Transaction] -> Ledger
simulateBank ledger ts = runST $ do
    b <- mkBank ledger
    mapM_ (transact b) ts
    readLedger b
```

Simulating a bank (contd.)

```
mkBank ledger = do
    let (n,as) = unzip ledger
    vs <- mapM newSTRef as
    return $ fromList $ zip n vs

transact b (Transaction cFrom cTo amt) = do
    modifySTRef' (b ! cFrom) (subtract amt)
    modifySTRef' (b ! cTo) (+ amt)

readLedger b = do
    let (n,v) = unzip $ toList b
    as <- mapM readSTRef v
    return $ zip n as
```

The ST Monad (continued)

- The type of `runST` is

```
runST :: (forall s. ST s a) -> a
```

This is so that references to mutable variables cannot escape a `runST` and the computation encapsulated by it really is pure.

- Neither

```
runST $ newSTRef 0
```

nor

```
\p -> runST $ readSTRef p
```

are accepted by the typechecker.

Mutable Vectors

- The `vector` package provides mutable vectors (`MVector`) which can be created, accessed and destructively modified in the `ST` monad (also variants for the `IO` monad).
- Like immutable vectors there are unboxed and storable variants.
- Conversion between mutable and immutable vectors is possible, but may copy the vector

```
freeze :: MVector#(S, A) -> ST#(S)(Vector#(A))  
thaw :: Vector#(A) -> ST#(S)(MVector#(S, A))
```

Sieve of Eratosthenes (again)

```
import qualified Data.Vector.Unboxed as VI
import qualified Data.Vector.Unboxed.Mutable as V
countPrimes :: Int -> Int
countPrimes n
    = VI.length $ VI.filter id $ mask
  where
    mask = VI.create mkMask
    mkMask :: ST s (V.MVector#(s, Bool))
    mkMask = do
      v <- V.replicate n True
      V.write v 0 False
      V.write v 1 False
      forM_ [2..(n-1)] $ \p -> do
        b <- V.read v p
        when b $
          forM_ [p*p, p*p+p..(n-1)] $ \i ->
            V.write v i False
    return v
```

REPA: REgular PArallel arrays

- Key data type

`Array r sh e`

- `r` can be `D`, `U` `F` and others.
- `sh` is a type of the shape class `Shape`. Example types

`Z :: Int :: Int`

Example values

`Z :: 0 :: 1`

- Moving from a delayed to manifest representation is under the programmer's control: `computeP`, `computeS`.

Parallelization for free

```
ubuntu@ip-172-31-40-69:~/examples-repa /usr/bin/time \
./dist/build/examples-repa/examples-repa \
80 20 Wave.jpg WaveEC2.jpg +RTS -N1
9.59user 0.11system 0:09.70elapsed 100%CPU
(0avgtext+0avgdata 556080maxresident)k
0inputs+43824outputs (0major+2447minor)pagefaults 0swaps

ubuntu@ip-172-31-40-69:~/examples-repa /usr/bin/time \
./dist/build/examples-repa/examples-repa \
80 20 Wave.jpg WaveEC2.jpg +RTS -N4
14.09user 0.14system 0:04.66elapsed 305%CPU
(0avgtext+0avgdata 559488maxresident)k
0inputs+43824outputs (0major+2910minor)pagefaults 0swaps
```